

# Homework 3

Young Won Kim (yk41) and Minh Pham (mnp7)

Spring 2017

## 1 MAP and MLE parameter estimation

### Part 1. Using MLE

$$D = \{x^{(i)} \mid 1 \leq i \leq m; x \in \{0, 1\}\}$$

Probability mass function of each  $x^{(i)}$ :

$$f(x^{(i)}; \theta) = \theta^{x^{(i)}} (1 - \theta)^{(1-x^{(i)})}$$

Likelihood function:

$$L(\theta) = \prod_{i=1}^m \theta^{x^{(i)}} (1 - \theta)^{(1-x^{(i)})}$$

Maximum likelihood function:

$$L(\theta) = \theta^{\sum_{i=1}^m x^{(i)}} (1 - \theta)^{(m - \sum_{i=1}^m x^{(i)})}$$

$$\theta = \operatorname{argmax}_{\theta} L(\theta) = \operatorname{argmax}_{\theta} \log(L(\theta))$$

$$\theta = \operatorname{argmax}_{\theta} \sum_{i=1}^m x^{(i)} \log(\theta) + (m - \sum_{i=1}^m x^{(i)}) \log(1 - \theta)$$

$$\frac{\partial \log(L(\theta))}{\partial \theta} = \frac{\sum_{i=1}^m x^{(i)}}{\theta} - \frac{(m - \sum_{i=1}^m x^{(i)})}{1 - \theta} \equiv 0$$

$$\sum_{i=1}^m x^{(i)} (1 - \theta) - (m - \sum_{i=1}^m x^{(i)}) \theta = 0$$

$$\sum_{i=1}^m x^{(i)} - m\theta = 0$$

$$\hat{\theta} = \frac{\sum_{i=1}^m x^{(i)}}{m}$$

## Part 2. Using MAP

$$P(\theta) = \theta^{(a-1)}(1-\theta)^{(b-1)}$$

$$\hat{\theta} = \operatorname{argmax}_{\theta} P(D | \theta) P(\theta) = \operatorname{argmax}_{\theta} \prod_{i=1}^m \theta^{x^{(i)}} (1-\theta)^{(1-x^{(i)})} \theta^{(a-1)} (1-\theta)^{(b-1)}$$

$$\hat{\theta} = \operatorname{argmax}_{\theta} \sum_{i=1}^m x^{(i)} \log(\theta) + (m - \sum_{i=1}^m x^{(i)}) \log(1-\theta) + \log(\theta)(a-1) + (b-1) \log(1-\theta)$$

$$\frac{\partial L(\theta)}{\partial \theta} = \frac{\sum_{i=1}^m x^{(i)}}{\theta} - \frac{(m - \sum_{i=1}^m x^{(i)})}{1-\theta} + \frac{a-1}{\theta} - \frac{b-1}{1-\theta} \equiv 0$$

$$0 = \sum_{i=1}^m x^{(i)}(1-\theta) - \theta(m - \sum_{i=1}^m x^{(i)}) + (a-1)(1-\theta) - \theta(b-1)$$

$$0 = \sum_{i=1}^m x^{(i)} + \theta(-m - a - b + 2) - 1 + a$$

With  $a = b = 1$ , we have:

$$\theta = \frac{\sum_{i=1}^m x^{(i)}}{m}$$

## 2 Logistic regression and Gaussian Naive Bayes

### Part 1 Logistic regression

$$P(y = 1 | x; \theta) = g(\theta^T x); \theta \in \mathbb{R}^{d+1}$$

$$P(y = 0 | x; \theta) = 1 - g(\theta^T x)$$

$$\text{where } g(\theta^T x) = \frac{1}{1 + e^{-\theta^T x}}$$

### Part 2 Gaussian Naive Bayes

$$P(y = 1 | x) = \frac{P(x|y=1)P(y=1)}{P(x)} = \frac{P(x|y=1)P(y=1)}{P(x|y=1)P(y=1) + P(x|y=0)P(y=0)}$$

$$P(y = 1 | x) = \frac{1}{1 + \frac{P(y=0)P(x|y=0)}{P(y=1)P(x|y=1)}}$$

$$P(y = 1 | x) = \frac{1}{1 + \exp(\ln \frac{P(y=0)P(x|y=0)}{P(y=1)P(x|y=1)})}$$

$$P(y = 1 | x) = \frac{1}{1 + \exp(\ln \frac{P(y=0)}{P(y=1)} + \ln \frac{P(x|y=0)}{P(x|y=1)})}$$

We have:

$$\frac{P(y=0)}{P(y=1)} = \frac{1-\phi}{\phi}$$

And:

$$\begin{aligned} \ln \frac{P(x|y=0)}{P(x|y=1)} &= \sum_{j=1}^d (\ln \frac{P(x_j|y=0)}{P(x_j|y=1)}) \\ \ln \frac{P(x|y=0)}{P(x|y=1)} &= \sum_{j=1}^d \ln \frac{\frac{1}{\sqrt{2\pi\sigma_j^2}} \exp(\frac{-(x_j+\mu_{j0})^2}{2\sigma_j^2})}{\frac{1}{\sqrt{2\pi\sigma_j^2}} \exp(\frac{-(x_j+\mu_{j1})^2}{2\sigma_j^2})} \\ \ln \frac{P(x|y=0)}{P(x|y=1)} &= \sum_{j=1}^d \ln \exp(\frac{(x_j-\mu_{j1})^2 - (x_j-\mu_{j0})^2}{2\sigma_j^2}) = \sum_{j=1}^d (\frac{\mu_{j0}-\mu_{j1}}{\sigma_j^2} x_j + \frac{\mu_{j1}^2-\mu_{j0}^2}{\sigma_j^2}) \end{aligned}$$

Therefore,

$$P(y = 1 \mid x) = \frac{1}{1 + \exp(w_0 + \sum_{j=1}^d w_j x_j)},$$

in which,

$$\begin{aligned} w_0 &= \ln \frac{1-\phi}{\phi} + \sum_{j=1}^d \frac{\mu_{j1}^2 - \mu_{j0}^2}{\sigma_j^2} \\ w_j &= \frac{\mu_{j0} - \mu_{j1}}{\sigma_j^2} \end{aligned}$$

$$P(y = 0 \mid x) = 1 - P(y = 1 \mid x) = \frac{\exp(w_0 + \sum_{j=1}^d w_j x_j)}{1 + \exp(w_0 + \sum_{j=1}^d w_j x_j)}$$

### Part 3 Uniform Class Priors

Since class 1 and class 0 are equally likely, we have:

$$\frac{P(y=0)}{P(y=1)} = \frac{1-\phi}{\phi} = 1$$

Therefore,

$$P(y = 1 \mid x) = \frac{1}{1 + \exp \sum_{j=1}^d (\frac{\mu_{j1}^2 - \mu_{j0}^2}{2\sigma_j^2} + \frac{\mu_{j0} - \mu_{j1}}{\sigma_j^2} x_j)}$$

We call:

$$\theta_0 = \frac{\mu_{j1}^2 - \mu_{j0}^2}{2\sigma_j^2}$$

$$\theta_j = \frac{\mu_{j0} - \mu_{j1}}{\sigma_j^2}$$

Therefore,

$$P(y = 1 \mid x) = \frac{1}{1 + \exp(\theta_0 + \sum_{j=1}^d (\theta_j x_j))} = \frac{1}{1 + \exp(\theta^T x)} \equiv P(y = 1 \mid x) \text{ for logistic regression}$$

### 3 Reject Option In Classifiers

#### Part 1 Minimum Risk

Risk for choosing an action  $i$  is:

$$R(\alpha_i | x) = \sum_{j=1}^C L(\alpha = i, y = j) P(y = j | x) \text{ for } i = 1, \dots, C$$

$$R(\alpha_i | x) = \sum_{j=1, j \neq i}^C \lambda_s P(y = j | x) = \lambda_s (1 - P(y = i | x))$$

Risk for a rejection is:

$$R(\alpha_{C+1} | x) = \lambda_r$$

Minimal risk that is obtained when we decide  $y = j$  is:

$$\lambda_r \leq \lambda_s (1 - P(y = i | x)) \text{ with } i = 1, \dots, C$$

Therefore,

$$P(y = i | x) \leq 1 - \frac{\lambda_r}{\lambda_s} \text{ with } i = 1, \dots, C$$

In addition, in an intuitive sense, we also have to choose the most probable class, which means  $P(y = j | x) \leq P(y = k | x)$  for all  $k$

#### Part 2 The loss ratio between reject and misclassification increases from 0 to 1

As  $\frac{\lambda_1}{\lambda_2}$  increases from 0 to 1, the loss of rejection increases. When  $\frac{\lambda_1}{\lambda_2} = 1$ , the loss of rejection is larger than the loss of misclassification, and we shouldn't choose the rejection option.

#### 4 One vs. all logistic regression and softmax regression

- As  $\theta$  is initially generated randomly, there is no preference for the class  $y$  belongs to, and therefore probability is equal for each classes,  $\frac{1}{10}$ . If we sum up across  $m$  examples and divide by  $m$ , the loss results to  $-\log(0.1)$ .

- Naive loss: 2.365103 computed in 26.325576s
- Vectorized loss: 2.365103 computed in 0.607100s
- Loss difference: 0.000000
- Gradient difference: 0.000000

```

one_vs_all on raw pixels final test set accuracy: 0.362000
[[465 59 21 24 19 35 26 60 201 90]
 [ 67 465 18 35 23 31 44 50 94 173]
 [123 65 193 77 96 89 151 89 69 48]
 [ 66 86 78 161 49 193 171 51 62 83]
 [ 65 38 103 64 234 90 194 128 36 48]
 [ 48 63 81 126 81 272 114 88 72 55]
 [ 31 53 67 102 85 78 457 52 29 46]
 [ 53 62 51 46 69 84 66 405 49 115]
 [143 78 8 25 9 34 22 20 547 114]
 [ 59 208 14 22 23 29 60 56 108 421]]

```

Figure 1: Confusion matrix of OVA

```

softmax on raw pixels final test set accuracy: 0.408100
[[484 51 50 23 14 24 25 45 203 81]
 [ 61 514 17 30 15 38 36 48 83 158]
 [ 96 59 258 68 119 86 164 65 58 27]
 [ 46 84 89 247 47 186 124 53 54 70]
 [ 63 38 110 67 293 76 179 110 31 33]
 [ 47 53 94 137 72 338 93 67 69 30]
 [ 18 53 51 98 102 80 496 38 25 39]
 [ 56 49 53 48 95 84 53 431 48 83]
 [161 74 8 16 5 51 10 14 550 111]
 [ 73 205 10 20 15 21 45 51 90 470]]

```

Figure 2: Confusion matrix of SOFTMAX

- The best value for learning rate and regularization strength are:  $lr=1.000000e-06$ ,  $reg=1.000000e+05$ . This resulted in validation accuracy of 41.6%.
- The best softmax classifier on the test set resulted in overall accuracy of 40.8%.
- Comparing OVA vs SOFTMAX on CIFAR-10: OVA yielded lower recall, precision, and accuracy compared to Softmax. Softmax works better on datasets where the features aren't obviously related to specific classes, which makes it a better classifier for CIFAR-10. Empirical data also supports this, as Softmax performed better than OVA in terms of recall, precision, and accuracy.
- The visualized coefficients corresponds to the importance of each pixel in determining the class. Bluer color represent higher correlation between the class and the pixel.

|                 | <b>OVA</b>    |                  | <b>Softmax</b> |                  |
|-----------------|---------------|------------------|----------------|------------------|
| <b>Label</b>    | <b>Recall</b> | <b>Precision</b> | <b>Recall</b>  | <b>Precision</b> |
| 0               | 0.465         | 0.4151786        | 0.484          | 0.438009         |
| 1               | 0.465         | 0.3950722        | 0.514          | 0.435593         |
| 2               | 0.193         | 0.3044164        | 0.258          | 0.348649         |
| 3               | 0.161         | 0.2360704        | 0.247          | 0.327586         |
| 4               | 0.234         | 0.3401163        | 0.293          | 0.377091         |
| 5               | 0.272         | 0.2909091        | 0.338          | 0.343496         |
| 6               | 0.457         | 0.3501916        | 0.496          | 0.404898         |
| 7               | 0.405         | 0.4054054        | 0.431          | 0.467462         |
| 8               | 0.547         | 0.4317285        | 0.55           | 0.45417          |
| 9               | 0.421         | 0.3528919        | 0.47           | 0.426497         |
| <b>Accuracy</b> | <b>0.362</b>  |                  | <b>0.4081</b>  |                  |

Figure 3: Comparison between OVA and SOFTMAX - precision, recall, accuracy

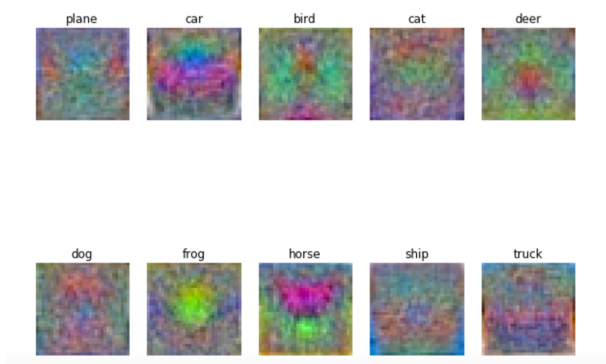


Figure 4: Visualization of the learned parameter matrix

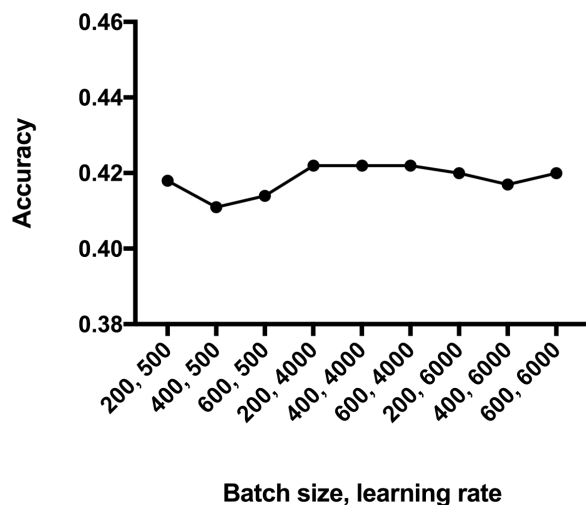


Figure 5: Other hyperparameters and their accuracies

### Extra Credit

Based on runs with different hyperparameters, we found that higher batch size and number of iterations tend to result in higher accuracy. However, increasing the batch size and number of iterations too much also leads to overfitting. The number of iterations that yielded the best accuracy was 4000, regardless of the batch size, highest accuracy being 0.422 (Fig.5).

Selected runs and their hyperparameters:

200 batch size, 500 iterations, lr  $5.0e-07$ , reg  $1.0e+05$ : train accuracy of 0.395, val accuracy of 0.418

400 batch size, 500 iterations, lr  $1.0e-06$ , reg  $1.0e+05$ : train accuracy of 0.403, val accuracy of 0.411

600 batch size, 500 iterations, lr  $1.0e-06$ , reg  $5.0e+05$ : train accuracy of 0.405, val accuracy of 0.414

200 batch size, 4000 iterations, lr  $5.0e-07$ , reg  $1.0e+05$ : train accuracy of 0.403, val accuracy of 0.422

400 batch size, 4000 iterations, lr  $1.0e-06$ , reg  $5.0e+04$ : train accuracy of 0.412, val accuracy of 0.422

600 batch size, 4000 iterations, lr  $1.0e-06$ , reg  $1.0e+05$ : train accuracy of 0.409, val accuracy of 0.422

200 batch size, 6000 iterations, lr  $5.0e-07$ , reg  $1.0e+05$ : train accuracy of 0.406, val accuracy of 0.420

400 batch size, 6000 iterations, lr  $5.0e-07$ , reg  $5.0e+04$ : train accuracy of 0.413, val accuracy of 0.417

600 batch size, 6000 iterations, lr  $5.0e-07$ , reg  $5.0e+04$ : train accuracy of 0.412, val accuracy of 0.420